

Introduction to the UNIX Operating System

Tutorial One

- Listing files and directories
- Making Directories
- Changing to a different Directory
- The directories . and ..
- Pathnames
- More about home directories and pathnames

Tutorial Two

- Copying Files
- Moving Files
- Removing Files and directories
- Displaying the contents of a file on the screen
- Searching the contents of a file

Tutorial Three

- Redirection
- Redirecting the Output
- Redirecting the Input
- Pipes

Tutorial Four

- Wildcards
- Filename Conventions
- Getting Help

Tutorial Five

- File system security (access rights)
- Changing access rights
- Processes and Jobs
- Listing suspended and background processes
- Killing a process

Tutorial Six

- Other Useful UNIX commands

1.1 Listing files and directories

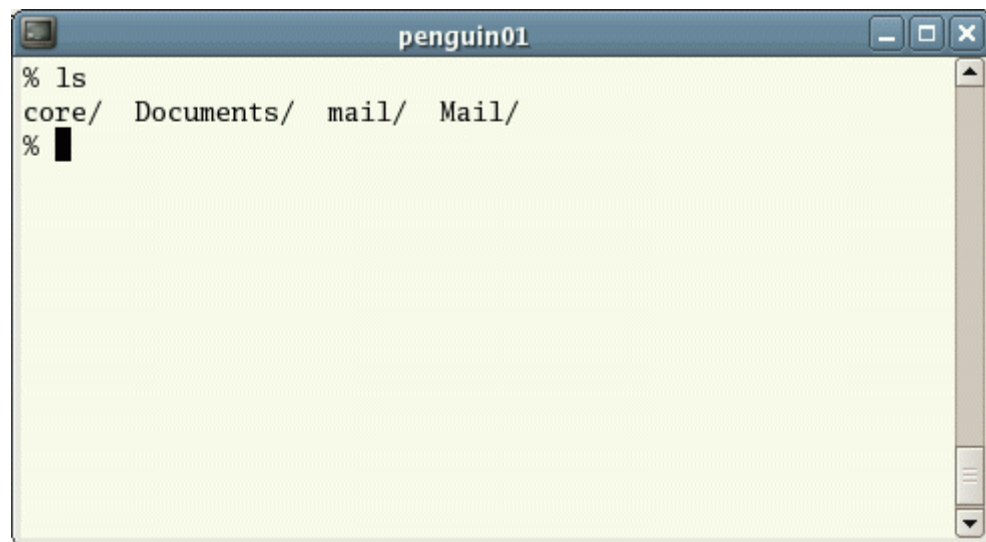
ls (list)

When you first login, your current working directory is your home directory. Your home directory has the same name as your user-name, for example, **jplante**, and it is where your personal files and subdirectories are saved.

To find out what is in your home directory, type

```
% ls
```

The **ls** command lists the contents of your current working directory.



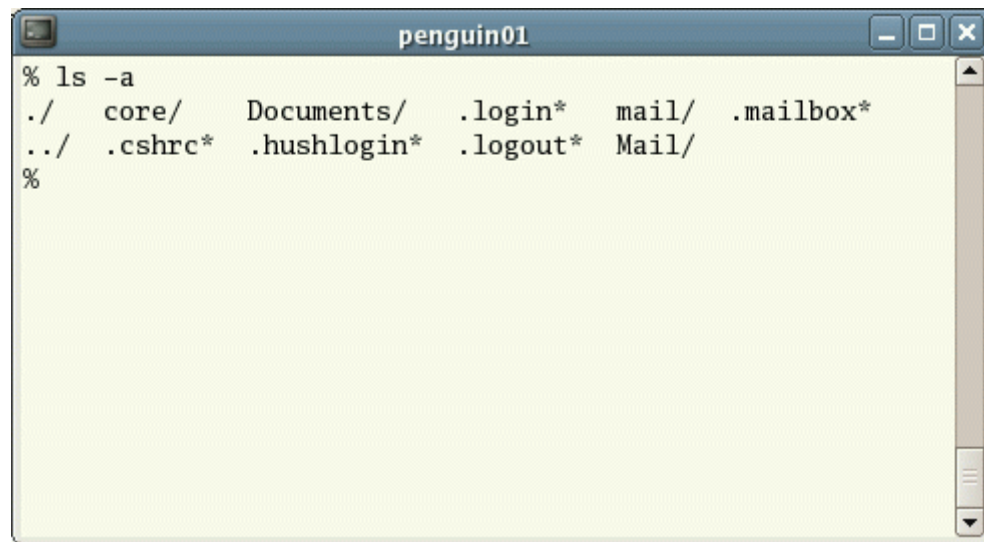
There may be no files visible in your home directory, in which case, the UNIX prompt will be returned. Alternatively, there may already be some files inserted by the System Administrator when your account was created.

ls does not, in fact, cause all the files in your home directory to be listed, but only those ones whose names do not begin with a dot (.). Files beginning with a dot (.) are known as hidden files and usually contain important program configuration information. They are hidden because you should not change them unless you are very familiar with UNIX!!!

To list *all* files in your home directory including those whose names begin with a dot, type

```
% ls -a
```

As you can see, `ls -a` lists files that are normally hidden.

A terminal window titled 'penguin01' with standard window controls (minimize, maximize, close). The prompt is '%'. The command 'ls -a' has been entered, and the output is displayed on two lines: './ core/ Documents/ .login* mail/ .mailbox*' and '../ .cshrc* .hushlogin* .logout* Mail/'. The prompt '%' is on the third line.

```
% ls -a
./ core/ Documents/ .login* mail/ .mailbox*
../ .cshrc* .hushlogin* .logout* Mail/
%
```

`ls` is an example of a command which can take options: `-a` is an example of an option. The options change the behavior of the command. There are online manual pages that tell you which options a particular command can take, and how each option modifies the behavior of the command. (See later in this tutorial)

1.2 Making Directories

mkdir (make directory)

We will now make a subdirectory in your home directory to hold the files you will be creating and using in the course of this tutorial. To make a subdirectory called **unixstuff** in your current working directory type

```
% mkdir unixstuff
```

To see the directory you have just created, type

```
% ls
```

1.3 Changing to a different directory

cd (change directory)

The command **cd *directory*** means change the current working directory to '*directory*'. The current working directory may be thought of as the directory you are in, i.e. your current position in the file-system tree.

To change to the directory you have just made, type

```
% cd unixstuff
```

Type **ls** to see the contents (which should be empty).

Exercise 1a

Make another directory inside the **unixstuff** directory called **backups**.

1.4 The directories . and ..

While you are still in the **unixstuff** directory, type

```
% ls -a
```

As you can see, in the **unixstuff** directory (and in all other directories), there are two special directories called **(.)** and **(..)**.

The current directory (.)

In UNIX, **(.)** means the current directory, so typing

```
% cd .
```

NOTE: there is a space between cd and the dot

means stay where you are (the **unixstuff** directory).

This may not seem very useful at first, but using **(.)** as the name of the current directory will save a lot of typing, as we shall see later in the tutorial.

The parent directory (..)

(..) means the parent of the current directory, so typing

```
% cd ..
```

will take you one directory up the hierarchy (back to your home directory). Try it now.

Note: typing **cd** with no argument always returns you to your home directory. This is very useful if you are lost in the file system.

1.5 Pathnames

pwd (print working directory)

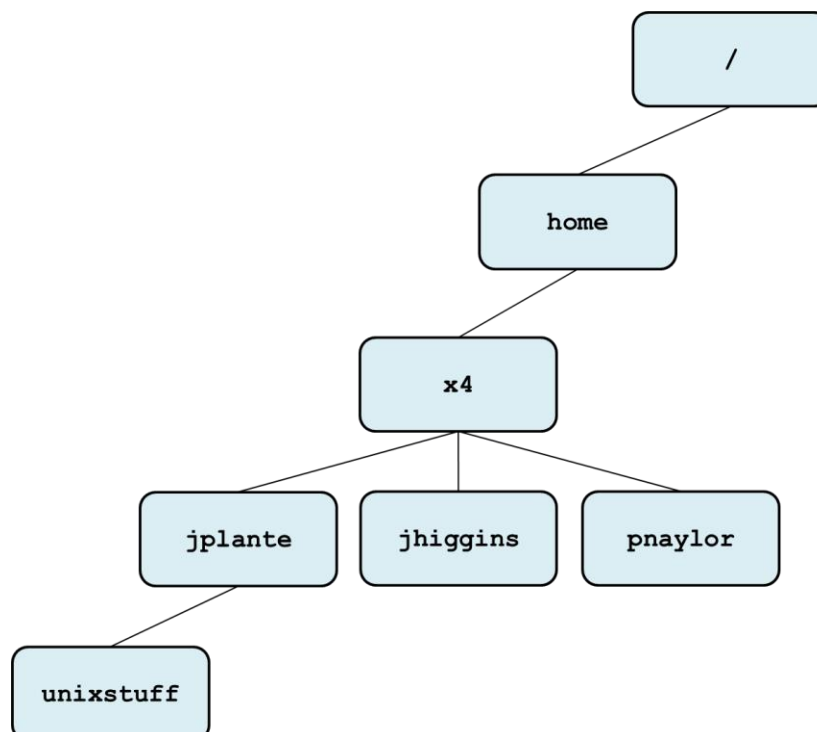
Pathnames enable you to work out where you are in relation to the whole file-system. For example, to find out the absolute pathname of your home-directory, type **cd** to get back to your home-directory and then type

```
% pwd
```

The full pathname will look something like this -

```
/home/x4/jplante
```

which means that **jplante** (your home directory) is in the sub-directory **x4** (the group directory), which in turn is located in the **home** sub-directory, which is in the top-level root directory called **" / "**.



Exercise 1b

Use the commands `cd`, `ls` and `pwd` to explore the file system.

(Remember, if you get lost, type `cd` by itself to return to your home-directory)

1.6 More about home directories and pathnames

Understanding pathnames

First type `cd` to get back to your home-directory, then type

```
% ls unixstuff
```

to list the contents of your **unixstuff** directory.

Now type

```
% ls backups
```

You will get a message like this -

```
ls: cannot access backups: No such file or directory
```

The reason is that **backups** is not in your current working directory. To use a command on a file (or directory) not in the current working directory (the directory you are currently in), you must either `cd` to the correct directory, or specify its full pathname. To list the contents of your **backups** directory, you must type

```
% ls unixstuff/backups
```

~ (your home directory)

Home directories can also be referred to by the tilde `~` character. It can be used to specify paths starting at your home directory. So typing

```
% ls ~/unixstuff
```

will list the contents of your **unixstuff** directory, no matter where you currently are in the file system.

What do you think the following command would list?

```
% ls ~
```

How about this one?

```
% ls ../../
```

Summary

Command	Meaning
ls	list files and directories
ls -a	list all files and directories
mkdir	make a directory
cd <i>directory</i>	change to named directory
cd	change to home-directory
cd ~	change to home-directory
cd ..	change to parent directory
pwd	display the path of the current directory

2.1 Copying Files

cp (copy)

cp file1 file2 is the command which makes a copy of **file1** in the current working directory and calls it **file2**.

What we are going to do now is to take a file stored in an open access area of the file system, and use the **cp** command to copy it to your **unixstuff** directory.

First, **cd** to your **unixstuff** directory.

```
% cd ~/unixstuff
```

Then at the UNIX prompt, type,

```
% wget http://www.cis.umassd.edu/~jplante/cis370/lab01/science.txt .
```

Note: Don't forget the dot **.** at the end. Remember, in UNIX, the dot means the current directory.

The above command means copy the file **science.txt** from a location online to the current directory, keeping the name the same. URLs are nothing more than file systems hosted on remote servers.

Exercise 2a

Create a backup of your **science.txt** file by copying it to a file called **science.bak**

2.2 Moving files

mv (move)

mv file1 file2 moves (or renames) **file1** to **file2**

To move a file from one place to another, use the **mv** command. This has the effect of moving rather than copying the file, so you end up with only one file rather than two.

It can also be used to rename a file, by moving the file to the same directory, but giving it a different name.

We are now going to move the file **science.bak** to your **backups** directory.

First, change directories to your **unixstuff** directory (can you remember how?). Then, inside the **unixstuff** directory, type

```
% mv science.bak backups/.
```

Type **ls** and **ls backups** to see if it has worked.

2.3 Removing files and directories

rm (remove), rmdir (remove directory)

To delete (remove) a file, use the **rm** command. As an example, we are going to create a copy of the **science.txt** file then delete it.

Inside your **unixstuff** directory, type

```
% cp science.txt tempfile.txt
% ls
% rm tempfile.txt
% ls
```

You can use the **rmdir** command to remove a directory (make sure it is empty first). Try to remove the **backups** directory. You will not be able to since UNIX will not let you remove a non-empty directory.

Exercise 2b

Create a directory called **tempstuff** using **mkdir** , then remove it using the **rmdir** command.

2.4 Displaying the contents of a file on the screen

clear (clear screen)

Before you start the next section, you may like to clear the terminal window of the previous commands so the output of the following commands can be clearly understood.

At the prompt, type

```
% clear
```

This will clear all text and leave you with the % prompt at the top of the window.

cat (concatenate)

The command cat can be used to display the contents of a file on the screen. Type:

```
% cat science.txt
```

As you can see, the file is longer than the size of the window, so it scrolls past making it unreadable.

less

The command less writes the contents of a file onto the screen a page at a time. Type

```
% less science.txt
```

Press the [**space-bar**] if you want to see another page, and type [**q**] if you want to quit reading. As you can see, **less** is used in preference to **cat** for long files.

head

The **head** command writes the first ten lines of a file to the screen.

First clear the screen then type

```
% head science.txt
```

Then type

```
% head -5 science.txt
```

What difference did the **-5** do to the **head** command?

tail

The **tail** command writes the last ten lines of a file to the screen.

Clear the screen and type

```
% tail science.txt
```

Critical Thinking: How can you view the last 15 lines of the file?

2.5 Searching the contents of a file

Simple searching using less

Using **less**, you can search through a text file for a keyword (pattern). For example, to search through **science.txt** for the word '**science**', type

```
% less science.txt
```

then, still in **less**, type a forward slash [/] followed by the word to search:

```
/science
```

As you can see, **less** finds and highlights the keyword. Type **[n]** to search for the next occurrence of the word.

grep (get rational expression)

grep is one of many standard UNIX utilities. It searches files for specified words or patterns.

First clear the screen, then type

```
% grep science science.txt
```

As you can see, **grep** has printed out each line containing the word **science**...

...Or has it?

Try typing

```
% grep Science science.txt
```

The **grep** command is case sensitive; it distinguishes between "Science" and "science."

To ignore upper/lower case distinctions, use the **-i** option:

```
% grep -i science science.txt
```

To search for a phrase or pattern, you must enclose it in single quotes (the apostrophe symbol). For example to search for "spinning top":

```
% grep -i 'spinning top' science.txt
```

Some of the other options of **grep** are:

- v** display those lines that do NOT match.
- n** precede each matching line with the line number.
- c** print only the total count of matched lines.

Try some of them and see the different results. Don't forget, you can use more than one option at a time. For example, the number of lines without the word "science" or "Science" is:

```
% grep -ivc science science.txt
```

wc (word count)

A handy little utility is the **wc** command, short for word count. To do a word count on **science.txt**, type

```
% wc -w science.txt
```

To find out how many lines the file has, type

```
% wc -l science.txt
```

Summary

Command	Meaning
cp <i>file1 file2</i>	copy file1 and call it file2
mv <i>file1 file2</i>	move or rename file1 to file2
rm <i>file</i>	remove a file
rmdir <i>directory</i>	remove a directory
cat <i>file</i>	display a file
less <i>file</i>	display a file a page at a time
head <i>file</i>	display the first few lines of a file
tail <i>file</i>	display the last few lines of a file
grep <i>'keyword' file</i>	search a file for keywords
wc <i>file</i>	count number of lines/words/characters in file

3.1 Redirection

Most processes initiated by UNIX commands write to the **standard output** (that is, they write to the terminal screen), and many take their input from the **standard input** (that is, they read it from the keyboard). There is also the **standard error**, where processes write their error messages, by default, to the terminal screen.

We have already seen one use of the **cat** command to write the contents of a file to the screen.

Now type **cat** without specifying a file to read

```
% cat
```

Then type a few words on the keyboard and press the **[Return]** key.

Finally hold the **[Ctrl]** key down and press **[d]** (written as **^D** for short) to end the input.

What has happened?

If you run the **cat** command without specifying a file to read, it reads the standard input (the keyboard), and on receiving the 'end of file' (**^D**), copies it to the standard output (the screen).

In UNIX, we can redirect both the input and the output of commands.

3.2 Redirecting the Output

We use the **>** symbol to redirect the output of a command. For example, to create a file called **list1** containing a list of fruit, type

```
% cat > list1
```

Then type in the names of some fruit. Press **[Return]** after each one.

```
pear  
banana  
apple  
^D {this means press [Ctrl] and [d] to stop}
```

What happens is the **cat** command reads the standard input (the keyboard) and the **>** redirects the output, which normally goes to the screen, into a file called **list1**.

To read the contents of the file, type

```
% cat list1
```

Exercise 3a

Using the above method, create another file called **list2** containing the following fruits: orange, plum, mango, grapefruit. Read the contents of **list2**.

3.2.1 Appending to a file

The form **>>** appends standard output to a file. So to add more items to the file **list1**, type

```
% cat >> list1
```

Then type in the names of more fruit:

```
peach
grape
orange
^D (Control D to stop)
```

To read the contents of the file, type

```
% cat list1
```

You should now have two files. One contains six fruit, the other contains four fruit.

We will now use the cat command to join (concatenate) **list1** and **list2** into a new file called **biglist**. Type

```
% cat list1 list2 > biglist
```

What this is doing is reading the contents of **list1** and **list2** in turn, and then outputting the text to the file **biglist**.

To read the contents of the new file, type

```
% cat biglist
```

3.3 Redirecting the Input

We use the **<** symbol to redirect the input of a command.

The command **sort** alphabetically or numerically sorts a list. Type

```
% sort
```

Then type in the names of some animals. Press **[Return]** after each one.

```
dog
cat
bird
ape
^D (control d to stop)
```

The output will be

```
ape
bird
cat
dog
```

Using **<** you can redirect the input to come from a file rather than the keyboard. For example, to sort the list of fruit, type

```
% sort < biglist
```

and now the sorted list will be output to the screen.

To output the sorted list to a file, type,

```
% sort < biglist > slist
```


Use **cat** to read the contents of the file **slist**.

3.4 Pipes

To see who is on the system with you, type

```
% who
```

One method to get a sorted list of names is to type,

```
% who > names.txt  
% sort < names.txt
```

This is a bit slow and you have to remember to remove the temporary file called **names.txt** when you have finished. What you really want to do is connect the output of the **who** command directly to the input of the sort command. This is exactly what pipes do. The symbol for a pipe is the vertical bar **|**.

For example, typing

```
% who | sort
```

will give the same result as above, but is quicker and cleaner.

To find out how many users are logged on, type

```
% who | wc -l
```

Exercise 3b

Using pipes, display all lines of **list1** and **list2** containing the letter 'p', and sort the result.

Summary

Command	Meaning
<code>command > file</code>	redirect standard output to a file

<code>command >> file</code>	append standard output to a file
<code>command < file</code>	redirect standard input from a file
<code>command1 command2</code>	pipe the output of command1 to the input of command2
<code>cat file1 file2 > file0</code>	concatenate file1 and file2 to file0
<code>sort</code>	sort data
<code>who</code>	list users currently logged in

4.1 Wildcards

The * wildcard

The character ***** is called a wildcard, and will match against zero or more characters in a file (or directory) name. For example, in your **unixstuff** directory, type

```
% ls list*
```

This will list all files in the current directory starting with **list....**

Try typing

```
% ls *list
```

This will list all files in the current directory ending with **....list**

The ? wildcard

The character **?** will match exactly one character.

So **?ouse** will match files like **house** and **mouse**, but not **spouse**.

Try typing

```
% ls ?list
```

4.2 Filename conventions

We should note here that a directory is merely a special type of file. So the rules and conventions for naming files apply also to directories.

In naming files, characters with special meanings such as `/ * & %` , should be avoided. Also, avoid using spaces within names. The safest way to name a file is to use only alphanumeric characters, that is, letters and numbers, together with `_` (underscore) and `.` (dot).

Good filenames	Bad filenames
project.txt	project
my_big_program.c	my big program.c
fred_dave.doc	fred & dave.doc

File names conventionally start with a lower-case letter, and may end with a dot followed by a group of letters indicating the contents of the file. For example, all files consisting of **C** code may be named with the ending **.c**. For example: **program1.c** . Then in order to list all files containing **C** code in your home directory, you need only type **ls *.c** in that directory.

4.3 Getting Help

Online Manuals

There are online manuals which give information about most commands. The manual pages tell you which options a particular command can take, and how each option modifies the behavior of the command. Type **man <command>** to read the manual page for a particular command.

For example, to find out more about the **wc** (word count) command, type

```
% man wc
```

Alternatively,

```
% whatis wc
```

gives a one-line description of the command, but omits any information about options, etc.

Apropos

When you are not sure of the exact name of a command,

```
% apropos keyword
```

will give you the commands with the given keyword in their manual page header. For example, try typing

```
% apropos copy
```

Summary

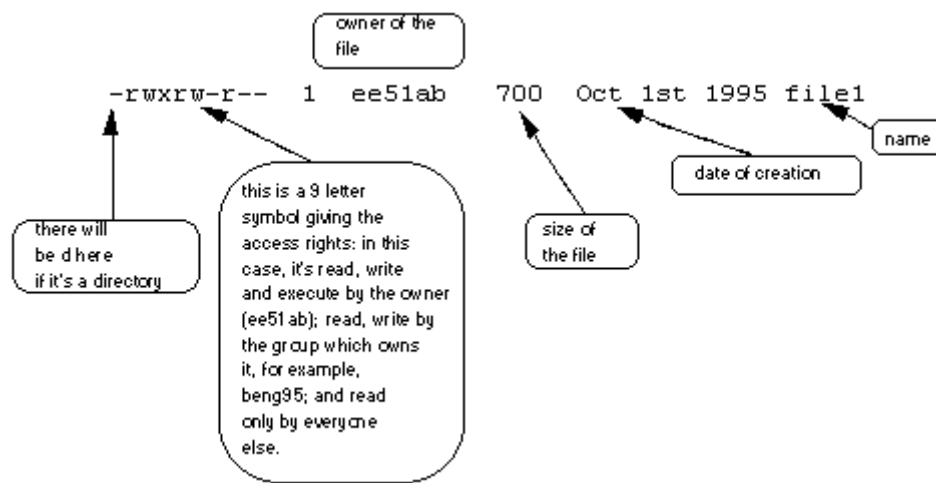
Command	Meaning
*	match any number of characters
?	match one character
man <i>command</i>	read the online manual page for a command
whatis <i>command</i>	brief description of a command
apropos <i>keyword</i>	match commands with keyword in their man pages

5.1 File system security (access rights)

In your **unixstuff** directory, type

```
% ls -l (l for long listing!)
```

You will see that you now get lots of details about the contents of your directory, similar to the example below.



Each file (and directory) has associated access rights, which may be found by typing `ls -l`. Also, `ls -lg` gives additional information as to which group owns the file (**cis370** in the following example):

```
-rwxrw-r-- 1 jplante cis370 2450 Sept29 11:52 file1
```

In the left-hand column is a 10 symbol string consisting of the symbols **d**, **r**, **w**, **x**, **-**, and, occasionally **s** or **S**. If **d** is present, it will be at the left hand end of the string, and indicates that the file is actually a directory: otherwise **-** will be the starting symbol of the string.

The 9 remaining symbols indicate the **file permissions**, or access rights, and are taken as three groups of 3 symbols each.

- The left group of 3 gives the file permissions for the **user** that owns the file (or directory) (**jplante** in the above example).
- The middle group of 3 gives the permissions for the **group** of people to whom the file (or directory) belongs (**cis370** in the above example).
- The rightmost group of 3 gives the permissions for all **others**.

The symbols **r**, **w**, etc., have slightly different meanings depending on whether they refer to a simple file or to a directory.

Access rights on files

- **r** (or **-**), indicates read permission (or otherwise); the presence or absence of permission to read/access and copy/download the file.
- **w** (or **-**), indicates write permission (or otherwise); the presence or absence of permission to change or overwrite/append to a file.
- **x** (or **-**), indicates execution permission (or otherwise); the presence or absence of permission to execute a file, for example if running a program executable (.exe) or object file (.o).

Access rights on directories

- **r** allows users to list files in the directory.
- **w** means that users may delete files from the directory or move files into it.
- **x** means the right to access files in the directory. This implies that you may read files in the directory provided you have read permission on the individual files.

So, in order to read a file, you must have execute permission on the directory containing that file, and hence on any directory containing that directory as a subdirectory, and so on, up the tree.

Some examples

<code>-rwxrwxrwx</code>	A file that everyone can read, write and execute (and delete).
<code>-rw-----</code>	A file that only the owner can read and write. Nobody else can read or write and nobody has execution rights (e.g. your mailbox file).

5.2 Changing access rights

chmod (changing a file mode)

Only the owner of a file can use **chmod** to change the permissions of a file. The options of **chmod** are as follows:

Symbol	Meaning
u	user
g	group
o	other
a	all
r	read
w	write (and delete)
x	execute (and access directory)
+	add permission
-	take away permission

For example, to remove read, write, and execute permissions on the file **biglist** for the group and others, type

```
% chmod go-rwx biglist
```

This will leave the other permissions unaffected.

To give read and write permissions on the file **biglist** to all (user, group, and others):

```
% chmod a+rw biglist
```

Exercise 5a

Try changing access permissions on the file **science.txt** and on the directory **backups**

Use **ls -l** to check that the permissions have changed.

BE CAREFUL: Don't lock yourself out of your own files!!!

5.3 Processes and Jobs

A process is an executing program identified by a unique PID (process identifier). To see information about your processes, with their associated PID and status, type

```
% ps
```

A process may be in the foreground, in the background, or be suspended. In general the shell does not return the UNIX prompt until the current process has finished executing.

Some processes take a long time to run and hold up the terminal. Backgrounding a long process has the effect that the UNIX prompt is returned immediately, and other tasks can be carried out while the original process continues executing.

Running background processes

To background a process, type an **&** at the end of the command line. For example, the command **sleep** waits a given number of seconds before continuing. Type

```
% sleep 10
```

This will wait 10 seconds before returning the command prompt **%**. Until the command prompt is returned, you can do nothing except wait.

To run **sleep** in the background, type

```
% sleep 10 &
```

```
[1] 6259
```

The **&** runs the job in the background and returns the prompt straight away, allowing you to run other programs while waiting for that one to finish.

The first line in the above example is typed in by the user; the next line, indicating job number and PID, is returned by the machine. The user will be notified of a job number (numbered from 1) enclosed in square brackets, together with a PID and is notified when a background process is finished. Backgrounding is useful for jobs which will take a long time to complete.

Backgrounding a current foreground process

At the prompt, type

```
% sleep 1000
```

You can suspend the process running in the foreground by typing **^Z**, i.e. hold down the [Ctrl] key and type [z]. Then to put it in the background, type

```
% bg
```

Note: do not background programs that require user interaction.

5.4 Listing suspended and background processes

When a process is running, backgrounded, or suspended, it will be entered into a list along with a job number. To examine this list, type

```
% jobs
```

An example of a job list could be

```
[1] Suspended sleep 1000
[2] Running netscape
[3] Running matlab
```

To restart (foreground) a suspended processes, type

```
% fg %jobnumber
```

For example, to restart sleep 1000, type

```
% fg %1
```

Typing **fg** with no job number will foreground the last suspended process.

5.5 Killing a process

kill (terminate or signal a process)

It is sometimes necessary to kill a process (for example, when an executing program is in an infinite loop).

To kill a job running in the foreground, type **^C** (**[Ctrl]** and **[c]**). For example, run

```
% sleep 100  
^C
```

To kill a suspended or background process, type

```
% kill %jobnumber
```

For example, run

```
% sleep 100 &  
% jobs
```

If it is job number 4, type

```
% kill %4
```

To check whether this has worked, examine the job list again to see if the process has been removed.

ps (process status)

Alternatively, processes can be killed by finding their process numbers (PIDs) and using kill *PID_number*

```
% sleep 1000 &  
% ps  
  
PID TT S TIME COMMAND  
20077 pts/5 S 0:05 sleep 1000  
21563 pts/5 T 0:00 netscape  
21873 pts/5 S 0:25 nedit
```

To kill off the process **sleep 1000**, type

```
% kill 20077
```

and then type **ps** again to see if it has been removed from the list.

If a process refuses to be killed, use the **-9** option, i.e. type

```
% kill -9 20077
```

Note: It is not possible to kill off other users' processes!!!

Summary

Command	Meaning
ls -lag	list access rights for all files
chmod [options] file	change access rights for named file
command &	run command in background
^C	kill the job running in the foreground
^Z	suspend the job running in the foreground
bg	background the suspended job
jobs	list current jobs
fg %1	foreground job number 1
kill %1	kill job number 1
ps	list current processes
kill 26152	kill process number 26152

6.1 Other useful UNIX commands

quota

All students are allocated a certain amount of disk space on the file system for their personal files, usually about 100Mb. If you go over your quota, you are given 7 days to remove excess files.

To check your current quota and how much of it you have used, type

```
% quota -v
```

df

The **df** command reports on the space left on the file system. For example, to find out how much space is left on the fileserver, type

```
% df .
```

du

The **du** command outputs the number of kilobytes used by each subdirectory. Useful if you have gone over quota and you want to find out which directory has the most files. In your home-directory, type

```
% du -s *
```

The **-s** flag will display only a summary (total size) and the ***** means all files and directories.

gzip

This reduces the size of a file, thus freeing valuable disk space. For example, type

```
% ls -l science.txt
```

and note the size of the file using **ls -l** . Then to compress science.txt, type

```
% gzip science.txt
```

This will compress the file and place it in a file called **science.txt.gz**

To see the change in size, type **ls -l** again.

To expand the file, use the **gunzip** command.

```
% gunzip science.txt.gz
```

zcat

zcat will read gzipped files without needing to uncompress them first.

```
% zcat science.txt.gz
```

If the text scrolls too fast for you, pipe the output through **less**.

```
% zcat science.txt.gz | less
```

file

file classifies the named files according to the type of data they contain, for example ASCII (text), pictures, compressed data, etc. To report on all files in your home directory, type

```
% file *
```

diff

This command compares the contents of two files and displays the differences. Suppose you have a file called **file1** and you edit some part of it and save it as **file2**. To see the differences type

```
% diff file1 file2
```

Lines beginning with a < denotes **file1**, while lines beginning with a > denotes **file2**.

find

This command searches through the directories for files and directories with a given name, date, size, or any other attribute you care to specify. It is a simple command but with many options. You can read the manual by typing **man find**.

To search for all files with the extension **.txt**, starting at the current directory (**.**) and working through all sub-directories, then printing the name of the file to the screen, type

```
% find . -name "*.txt" -print
```

To find files over 1Mb in size, and display the result as a long listing, type

```
% find . -size +1M -ls
```

history

The C shell keeps an ordered list of all the commands that you have entered. Each command is given a number according to the order it was entered.

```
% history (show command history list)
```

If you are using the C shell, you can use the exclamation character (**!**) to recall commands easily.

```
% !! (recall last command)
```

```
% !-3 (recall third most recent command)
```

```
% !5 (recall 5th command in list)
```

```
% !grep (recall last command starting with grep)
```

You can also reenter the last command you typed by hitting the up arrow on your keyboard when at the command line. You can navigate through previous commands with the up and down arrows.